

HPC Tools 2024, Exercise 9 solutions

Bruno Oliveira Cattelan, Laurent Chôné

Problem 1

OpenMP as a tool tries to allow for parallelization with minimal impact in the code. For example, both serial and parallel versions used in this exercise used the same code. The serial version is achieved by not enabling OpenMP at compile time.

```

1  // *** C++ ***
2  // Simulation proper
3
4  for (n=0;n<maxt;n++) {
5      #pragma omp parallel for simd
6          for (i=0;i<nat;i++) v0[i]=v[i];
7
8      #pragma omp parallel for
9          for (i=0;i<nat;i++)
10             // New potential energy and acceleration
11             accel(nat,i,&ep[i],&a[i],box,x);
12
13     #pragma omp parallel for private(vave) firstprivate(ek, v, x, a, dt, box)
14         for (i=0;i<nat;i++) {
15
16             // Leap frog integration algorithm: update position and velocity
17             v[i]=v[i]+dt*a[i];
18             x[i]=x[i]+dt*v[i];
19
20             // Check periodic boundary conditions
21             if (x[i]<0.0 ) x[i]=x[i]+box;
22             if (x[i]>=box) x[i]=x[i]-box;
23
24             // Calculate kinetic energy (note: mass=1)
25             vave=(v0[i]+v[i])/2.0;
26             ek[i]=1.0/2.0*vave*vave;
27
28         }
29
30
31     // Calculate and print total potential end kinetic energies
32     // and their sum that should be conserved.
33     epsum=eksum=0.0;
34     #pragma omp parallel for reduction(+:epsum)
35         for (i=0;i<nat;i++) epsum+=ep[i];
36     #pragma omp parallel for reduction(+:eksum)
37         for (i=0;i<nat;i++) eksum+=ek[i];
38
39 }

```

From Figure 1 we see that although the speedup is not ideal, the result is actually much faster with 16 threads than the serial counterpart. For this experiment we used as input the following sequence 5000000.001100011000. Meaning 500000 atoms, 0.001 as the time interval, 1000 timesteps and velocities initialized with 1. The last two numbers refer to the frequency of the outputs of the simulation. We leave them here to ensure the code is not optimized away by our *O2* optimization level. We also see that for larger than 16 threads there is a drop in our parallelization efficiency, which seems to pick back up after. This comes from the fact that the computer used for this had indeed 16 physical processors. If we continued the experiments for even larger number of threads, we would see this behavior repeated in intervals of 16.




Figure 1: OpenMP vs serial

```

1  !*** Fortran ***
2  ! Simulation proper
3  time_loop: do n=1,maxt
4
5      epsum = 0.0
6      eksum = 0.0
7
8      v0=v
9
10     !$omp parallel do
11     atom_loop1: do i=1,nat
12         ! New potential energy and acceleration
13         call accel(i,ep(i),a(i))
14     end do atom_loop1
15     !$omp end parallel do
16
17     !$omp parallel do
18     atom_loop2: do i=1,nat
19         ! Leap frog integration algorithm: update position and velocity
20         v(i)=v(i)+dt*a(i)
21         x(i)=x(i)+dt*v(i)
22         ! Check periodic boundary conditions
23         if (x(i)<0.0 ) x(i)=x(i)+box
24         if (x(i)>=box) x(i)=x(i)-box
25         ! Calculate kinetic energy (note: mass=1)
26         ek(i)=1.0/2.0*((v0(i)+v(i))/2.0)**2
27     end do atom_loop2
28     !$omp end parallel do
29
30     ! Calculate and print total potential end kinetic energies
31     ! and their sum that should be conserved.
32     !$omp parallel do reduction(+:epsum)
33     epsum_red: do i=1,nat
34         epsum=epsum+ep(i)
35     end do epsum_red
36     !$omp end parallel do
37     !$omp parallel do reduction(+:eksum)
38     eksum_red: do i=1,nat
39         eksum=eksum+ek(i)
40     end do eksum_red
41     !$omp end parallel do

```

```
42 end do time_loop
```

Finally, it is important to keep in mind Amdahl's law. The Figure 2 shows the theoretical speedup expected for different parallelizable portions of a code. However, super-linear speedups can also be achieved in specific cases!




Figure 2: Amdahl's law (Wikipedia)